



## Working with kafka

School of Science and Engineering

Nouamane Zanboui

Meriem Lmoubariki

Adam Sterheltou

Dr. Tajjeeddine Rachidi

Al Akhawayn University in Ifrane

November 11/12/2024

## **Table of Contents**

<b>Introduction .....</b>	<b>1</b>
<b>Question 1 .....</b>	<b>2</b>
<b>Question 2.....</b>	<b>3</b>
<b>Question 3.....</b>	<b>4</b>
<b>Question 4.....</b>	<b>5</b>
<b>Question 5 .....</b>	<b>6</b>
<b>Challenges.....</b>	<b>7</b>
<b>References .....</b>	<b>8</b>

## **Introduction**

In this report, we explore how Kafka, a powerful tool for handling real-time data, can be set up and used to manage streams of log entries. Kafka is widely used because it allows large amounts of data to be processed quickly and efficiently. This makes it invaluable for businesses that need to analyze data in real time, such as monitoring website traffic through log files.

The purpose of this assignment is to understand the fundamental roles of Kafka as a message broker, compare it with other similar technologies, and practically implement it by streaming data from log files. This hands-on experience will help us see how Kafka fits into the bigger picture of data management and processing.

### **Question 1 : Explain the Function of Kafka as a Message Broker**

Apache Kafka is a distributed system that serves as a message broker, enabling the fast, reliable exchange of data across various applications. As a message broker, Kafka allows different applications to communicate by sending, storing, and retrieving data, all through a central hub. Rather than connecting directly, systems send their data to Kafka, which then organizes stores, and delivers it to other systems that need this information (Kafka Documentation, 2023).

Kafka uses a **publish-subscribe model** for managing data. In this model, **producers** (systems that generate data) send messages to specific topics in Kafka. Topics act like channels that organize data by category, making it easy for **consumers** (systems or applications that need data) to subscribe to the topics they need. Kafka's durable storage system saves messages to disk, ensuring consumers can access data even if they were offline when the messages were initially sent. This setup is useful for a wide range of applications, including tracking website activity, managing logs, and processing real-time financial transactions.

Kafka is highly efficient and scalable, capable of handling large data volumes with low delay. Its ability to decouple producers from consumers also makes it flexible for complex, data-driven systems, where applications need to operate independently without direct dependency on each other (Kreps et al., 2011; Confluent, 2023).

## Question 2 : Comparison of Apache Kafka with Other Message Brokers and Pub/Sub Systems

(Pivotal Software, 2023; Red Hat, 2023)

Criteria	Apache Kafka	RabbitMQ	ActiveMQ	AWS SNS (Simple Notification Service)
<b>Architecture</b>	Distributed, log-based	Centralized, queue-based	Centralized, queue-based	Managed pub/sub service
<b>Message Delivery</b>	At least once, supports exactly-once	At most once, supports at-least-once	At most once, supports at-least-once	At least once
<b>Persistence</b>	Highly durable, persists messages on disk	Optional, mainly in-memory	Optional, persistent and in-memory	No built-in persistence
<b>Scalability</b>	Highly scalable, supports partitioning	Limited horizontal scaling	Limited horizontal scaling	Scalable, managed by AWS
<b>Throughput</b>	Very high throughput, optimized for logs	Moderate, better for transactional data	Moderate	High, but depends on AWS infrastructure
<b>Latency</b>	Low latency, ideal for real-time processing	Low latency, but may vary	Higher latency than Kafka	Low latency, managed by AWS
<b>Ordering</b>	Guarantees ordering per partition	Ordering optional, but not guaranteed	Supports ordering with configuration	No strict ordering
<b>Protocol</b>	Custom binary protocol, efficient (TCP)	AMQP, HTTP	OpenWire, STOMP, AMQP	HTTP-based
<b>Use Cases</b>	High-throughput logging, event streaming	Real-time updates, task queues	Enterprise messaging, legacy systems	Notifications, broadcast messaging
<b>Durability</b>	Very durable storage, keeps logs	Less durable, focuses on quick delivery	Moderate durability, disk or memory	Managed durability by AWS
<b>Management</b>	Requires setup, managed solutions exist	Requires setup, easier to manage	Requires setup, suited for smaller setups	Fully managed by AWS

**Apache Kafka:** Kafka is well-suited for applications that handle a lot of data at high speeds, like log processing, real-time analytics, and data pipelines. Its architecture is distributed, meaning it can be spread across multiple servers, making it very scalable. Kafka stores messages on disk, so even if the system goes down, data is preserved. It's best for large-scale applications where data needs to be processed quickly and reliably.

**RabbitMQ:** RabbitMQ is a popular choice for tasks that need reliable message delivery but don't require the massive scale that Kafka supports. It's easier to set up and integrates well with many messaging protocols like AMQP and HTTP. RabbitMQ is commonly used for smaller tasks like job queues or website updates, where it's important that messages are delivered reliably but not necessarily at the extremely high speeds Kafka can handle.

**ActiveMQ:** ActiveMQ is often used in businesses that need reliable messaging but don't require the high throughput that Kafka offers. It works well with older (or "legacy") systems and supports a range of messaging protocols. ActiveMQ can persist messages if needed, but it is typically used in smaller applications that need straightforward message passing rather than massive data streaming.

**AWS SNS:** Amazon Simple Notification Service (SNS) is a fully managed service provided by AWS, meaning Amazon handles setup and scaling. It's ideal for sending notifications and alerts, especially for applications already running in the AWS environment. Unlike Kafka, AWS SNS does not keep messages after they're sent, making it less suited for applications that need to retain message history. It's perfect for cases where instant notifications are needed, such as mobile alerts or system alerts.

**Scalability and Throughput:** Kafka is designed to handle extremely high data loads, making it ideal for big data applications where large amounts of data need to be processed in real-time. RabbitMQ and ActiveMQ don't scale to the same degree, but they are great for handling smaller tasks that need reliable message delivery. AWS SNS is highly scalable within the AWS ecosystem, making it convenient for applications hosted on AWS.

**Persistence and Durability:** Kafka is highly durable, storing messages on disk so they can be retrieved even after downtime or disruptions. This makes it useful for applications that need data stored reliably. In contrast, RabbitMQ and ActiveMQ offer optional message storage, while AWS SNS doesn't keep messages once they're sent.

**Ideal Use Cases:** Kafka is the go-to choice for high-volume, real-time applications like log processing and streaming analytics. RabbitMQ and ActiveMQ are preferred in scenarios needing reliable, smaller-scale messaging, and AWS SNS works well for quick notifications and mobile push alerts.

### Question 3

In this part, we explain how we set up and got Apache Kafka running using Docker. Our main aim was to create a working system that could handle and manage streams of log data efficiently. We chose Docker because it makes the setup process easier and keeps everything neat by running Kafka and Zookeeper in separate, contained environments.

We chose the Confluent Kafka Docker image because it's developed by the creators of Kafka, ensuring it's highly reliable and integrates smoothly with necessary components like Zookeeper. Confluent provides a user-friendly setup with robust support and documentation, making it an ideal choice for both ease of use and access to advanced features for effective learning and deployment in our project (Confluent, 2023).

#### Step 1:

##### Setting Up the Environment

- **Docker Environment:** We began by ensuring Docker was installed on our system. (It was obvious that it was installed from the beginning of the semester)

#### Step 2:

##### Acquiring Necessary Docker Images

- **Downloading Images:** We executed the following commands to download the latest Docker images for Zookeeper and Kafka from Confluent. This step ensures that we have reliable and up-to-date versions of both services, which are essential for Kafka's operation.

##### Commands used:

```
docker pull confluentinc/cp-zookeeper
```

```
docker pull confluentinc/cp-kafka
```

#### Step 3:

##### Launching Zookeeper

- **Running Zookeeper Container:** We used the following command to start a Zookeeper container:

```
docker run --name zookeeper -p 2181:2181 -e ZOOKEEPER_CLIENT_PORT=2181 confluentinc/cp-zookeeper
```

This command sets up a Zookeeper instance necessary for Kafka's management, specifying the client port to enable Kafka's connection to it. The -p option maps the port from the container to our host.

## Step 4:

### Deploying Kafka Broker

- **Starting Kafka Container:** To deploy Kafka, we issued:

```
docker run --name kafka-broker -p 9092:9092 --link zookeeper:zookeeper -e KAFKA_ZOOKEEPER_CONNECT=zookeeper:2181 -e KAFKA_ADVERTISED_LISTENERS=PLAINTEXT://localhost:9092 -e KAFKA_OFFSETS_TOPIC_REPLICATION_FACTOR=1 confluentinc/cp-kafka
```

This command initializes a Kafka broker, linking it to the Zookeeper container. It configures Kafka to use Zookeeper for coordination and sets Kafka to advertise itself on localhost for connections. The environment variables specify configuration settings critical for Kafka's interaction with Zookeeper and clients.

## Step 5:

### Testing the Installation

- **Creating and using a Kafka Topic:**

- We created a topic named "test-topic" using:

```
docker exec kafka-broker kafka-topics --create --topic test-topic --bootstrap-server localhost:9092 --partitions 1 --replication-factor 1
```

This command establishes a topic where messages will be stored. Setting partitions and replication factors defines how data is distributed and duplicated across the cluster.

To test the setup, we used the Kafka console producer and consumer to send and receive messages:

```
docker exec -it kafka-broker kafka-console-producer --topic test-topic --bootstrap-server localhost:9092
```

```
docker exec -it kafka-broker kafka-console-consumer --topic test-topic --from-beginning --bootstrap-server localhost:9092
```

### Results and Verification

```

Windows PowerShell x Windows PowerShell x Windows PowerShell x Windows PowerShell x + Windows PowerShell
Windows PowerShell
copyright (C) Microsoft Corporation. All rights reserved.

Install the latest PowerShell for new features and improvements! https://aka.ms/PSWindows

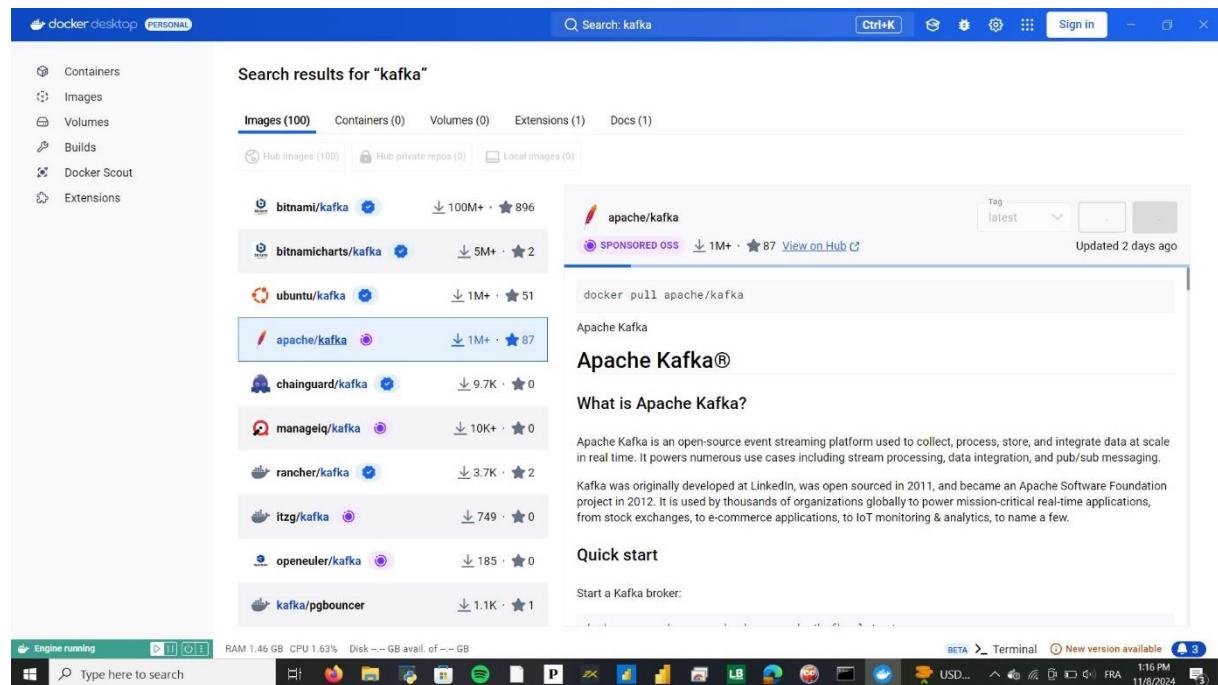
PS C:\Users\zambo> docker images
REPOSITORY          TAG      IMAGE ID      CREATED        SIZE
bitnami/zookeeper   latest   4412b94e1b99  48 hours ago   931MB
bitnami/kafka       latest   f50d4a0404f0  48 hours ago   1.07GB
cassandra          latest   5d795c114091  3 weeks ago    579MB
apache/kafka        latest   22c4baa38873  3 Weeks ago    599MB
apache/cassandra   4.0.1    1313a0a0a000  3 Weeks ago    3.07GB
bitnami/spark       latest   87a21bd1aa61  6 weeks ago    2.82GB
confluentinc/cp-kafka
latest   653f49c9c1cf  2 months ago   1.39GB
confluentinc/cp-zookeeper
latest   6d90a1a0a000  2 months ago   1.39GB
zookeeper           3.7     6c5ec2fa02ab  3 months ago   427MB
dixitf/image-tools  1.0.2    2514bf8e0597  7 months ago   19.7MB
mysql/mysql-server  latest   93a93a0a0000  9 months ago   4.09GB
louwu/hadoop-hive   latest   cb969afef439  3 years ago    3.6GB
bde2028/hadoop-namenode
latest   fd7f711b0801  4 years ago    2.05GB
bde2028/hadoop-mapreduce
latest   07aef0a0a000  4 years ago    2.05GB
bde2028/hadoop-datanode
latest   354899bcbe9f  4 years ago    2.05GB
harishkhan/hbase    latest   c65ce56799f5  4 years ago    766MB
bde2028/hadoop-mapreduce
latest   3.0.0-hadoop3.2.1-javab  91ca2a7fe8616  4 years ago    2.05GB
bde2028/hadoop-namenode
latest   2.0.0-hadoop3.2.1-javab  51ad9293ec52  4 years ago    2.05GB
bde2028/hadoop-zkserver
latest   2.0.0-hadoop3.2.1-javab  91ca2a7fe8616  4 years ago    2.05GB
bde2028/hadoop-datanode
latest   ddf4ec9ad55af  4 years ago    2.05GB
vishnukant/cassandra-oc-anyuser
latest   9575b3e709d  5 years ago    480MB
bde2028/hadoop-phoenix
latest   4ae8a31e062ed  6 years ago    1.94GB
bde2028/hadoop-phoenix
latest   ba849ed538d9  7 years ago    500MB
PS C:\Users\zambo> docker pull confluentinc/cp-zookeeper
docker pull confluentinc/cp-zookeeper
latest: Pulling from confluentinc/cp-zookeeper
Digest: sha256:86ffef47c8b855298bb21c156e99a3859df7a1de46379e43a137d32e8ea3
status: Image is up to date for confluentinc/cp-zookeeper:latest
Status: Downloaded newer image for confluentinc/cp-zookeeper:latest

What's next:
  View a summary of image vulnerabilities and recommendations + docker scout quickview confluentinc/cp-zookeeper
PS C:\Users\zambo> docker run -d --name zookeeper -p 2181:2181 --ZOOKEEPER_CLIENT_PORT=2181 confluentinc/cp-zookeeper
f6cc6101ea0dc4670bed029746b396cfa0d3952b59eca141d6e67183096a
PS C:\Users\zambo> docker run -d --name kafka-broker -p 9092:9092 --KAFKA_ZOOKEEPER_CONNECT=zookeeper:2181 --KAFKA_ADVERTISED_LISTENERS=PLAINTEXT://localhost:9092 --KAFKA_OFFSETS_TOPIC_REPLICATION_FACTOR=1 --link zookeeper:zookeeper
e3b3c3a65cc7af0b2f21806fb1d408893fed45119b3decdd7f719b883627
PS C:\Users\zambo> docker exec -it kafka-broker kafka-topics --create --topic test-topic --bootstrap-server localhost:9092 --partitions 1 --replication-factor 1
Created topic test-topic.

What's next:
  Try Docker Debug for seamless, persistent debugging tools in any container or image + docker debug kafka-broker
  Learn more at https://docs.docker.com/go/debug-cl/
  PS C:\Users\zambo> docker run -d --name kafka-console-producer --topic test-topic --bootstrap-server localhost:9092
>Hello this is an assignment of big data of Kafka.
>c
What's next:
  Try Docker Debug for seamless, persistent debugging tools in any container or image + docker debug kafka-broker
  Learn more at https://docs.docker.com/go/debug-cl/
  PS C:\Users\zambo> docker exec -it kafka-broker kafka-console-consumer --topic test-topic --bootstrap-server localhost:9092 --from-beginning
>Hello this is an assignment of big data of Kafka.
|
```

The attached previous screenshot (Figure 1) shows the successful execution of these commands and the message flow verification in Kafka. This visual confirmation is crucial for demonstrating the Kafka setup's functionality.

In the following figure, you will find another image with is bitnami KAFKA, it also worked.



```
Command Prompt - docker exec -it cswork-kafka-1 kafka-console-producer.sh --broker-list localhost:9092 --topic test
Microsoft Windows [Version 10.0.19045.4170]
(c) Microsoft Corporation. All rights reserved.

C:\Users\LENOVO>java -version
java version "21.0.1" 2023-10-17 LTS
Java(TM) SE Runtime Environment (build 21.0.1+12-LTS-29)
Java HotSpot(TM) 64-Bit Server VM (build 21.0.1+12-LTS-29, mixed mode, sharing)

C:\Users\LENOVO>docker pull apache/kafka
Using default tag: latest
latest: Pulling from apache/kafka
Digest: sha256:22c4bea38875408e8fe52aca8e3a6ee67f9aa0090db59af99a2f6647558db5
Status: Downloaded newer image for apache/kafka:latest
docker.io/apache/kafka:latest

What's next:
  View a summary of image vulnerabilities and recommendations → docker scout quickview apache/kafka

C:\Users\LENOVO>cd C:\cs work

C:\cs work>docker-compose up -d
time="2024-11-08T13:18:43+01:00" level=warning msg="C:\\cs work\\docker-compose.yml: the attribute `version` is obsolete, it will be ignored, please remove it to avoid potential confusion"
[*] Running 2/35
  □ zookeeper Pulled
    e057c74597c7 Download complete 69.3s
    e057c74597c7 Download complete 3.4s
    3d3a7dd3a3b1 Download complete 41.4s
    6b9611650a73 Download complete 3.4s
    666c214f6385 Download complete 3.4s
    a3ed95caeb02 Download complete 3.4s
    8b130a9baa49 Download complete 54.7s
    76eea4448d9b Download complete 3.4s
    8b666998876c6 Download complete 3.4s
    ef38b711a50f Download complete 3.4s

Windows PowerShell                                     Result   10:39 PM  11/8/2024
Type here to search
```

```
Command Prompt - docker exec -it cswork-kafka-1 kafka-console-producer.sh --broker-list localhost:9092 --topic test
3.4s
  □ kafka Pulled
    44b062e78fd7 Download complete 63.6s
    44b062e78fd7 Download complete 5.7s
    10c9a58bd495 Download complete 31.4s
    42c077c1079b Download complete 55.3s
    539ec416bc55 Download complete 5.3s
    03346d650161 Download complete 52.0s
    b3ba9647f279 Download complete 5.3s
    ed9bd501c190 Download complete 3.3s
[*] Running 3/3
  □ Network cswork_default Created
  □ Container cswork-kafka-1 Started
  □ Container cswork-zookeeper-1 Started

C:\cs work>docker exec -it cswork-kafka-1 kafka-topics.sh --create --topic test --bootstrap-server localhost:9092 --partitions 1 --replication-factor 1
Created topic test.

What's next:
  Try Docker Debug for seamless, persistent debugging tools in any container or image → docker debug cswork-kafka-1
  Learn more at https://docs.docker.com/go/debug-cli/

C:\cs work>docker exec -it cswork-kafka-1 kafka-console-producer.sh --broker-list localhost:9092 --topic test
>docker exec -it cswork-kafka-1 kafka-console-consumer.sh --bootstrap-server localhost:9092 --topic test --from-beginning
>
```

We did identically the same process but only with a different image.

docker desktop PERSONAL

Containers [Give feedback](#)

Container CPU usage: 0.62% / 800% (8 CPUs available)

Container memory usage: 501.07MB / 15.12GB

Show charts

Search  Only show running containers

	Name	Container ID	Image	Port(s)	CPU (%)	Last started	Actions
○	cassandraregionserver2	8177903884c2	cassandra:latest	-	0%	5 days ago	▷ ⚡ 🗑
○	cassandraregionserver3	303961f7984e	cassandra:latest	-	0%	5 days ago	▷ ⚡ 🗑
○	cassandraregionserver1	a81023374e3d	cassandra:latest	-	0%	5 days ago	▷ ⚡ 🗑
○	cass_cluster	0935fc8455bf	cassandra:latest	-	0%	5 days ago	▷ ⚡ 🗑
○	practical_buck	d80c47b4ac7b	jupyter/pyspark-notebook	8888:8888	0%	1 month ago	▷ ⚡ 🗑
○	hadoop-hive-container	e29e5b5fd20e	loum/hadoop-hive:<none>	10000:10000...	0%	1 month ago	▷ ⚡ 🗑
○	cswork	-	-	-	0.6%	8 minutes ago	▀ ⚡ 🗑
●	kafka-1	efd36383724b	wurstmeister/kafka:<none>	9092:9092	0.51%	8 minutes ago	▀ ⚡ 🗑
●	zookeeper-1	ab20e5a28a07	wurstmeister/zookeeper	2181:2181	0.09%	8 minutes ago	▀ ⚡ 🗑

Showing 9 items

Engine running RAM 1.59 GB CPU 0.13% Disk -- GB avail. of -- GB

Type here to search

BETA Terminal New version available 10:14 PM 11/11/2024

#### Question 4

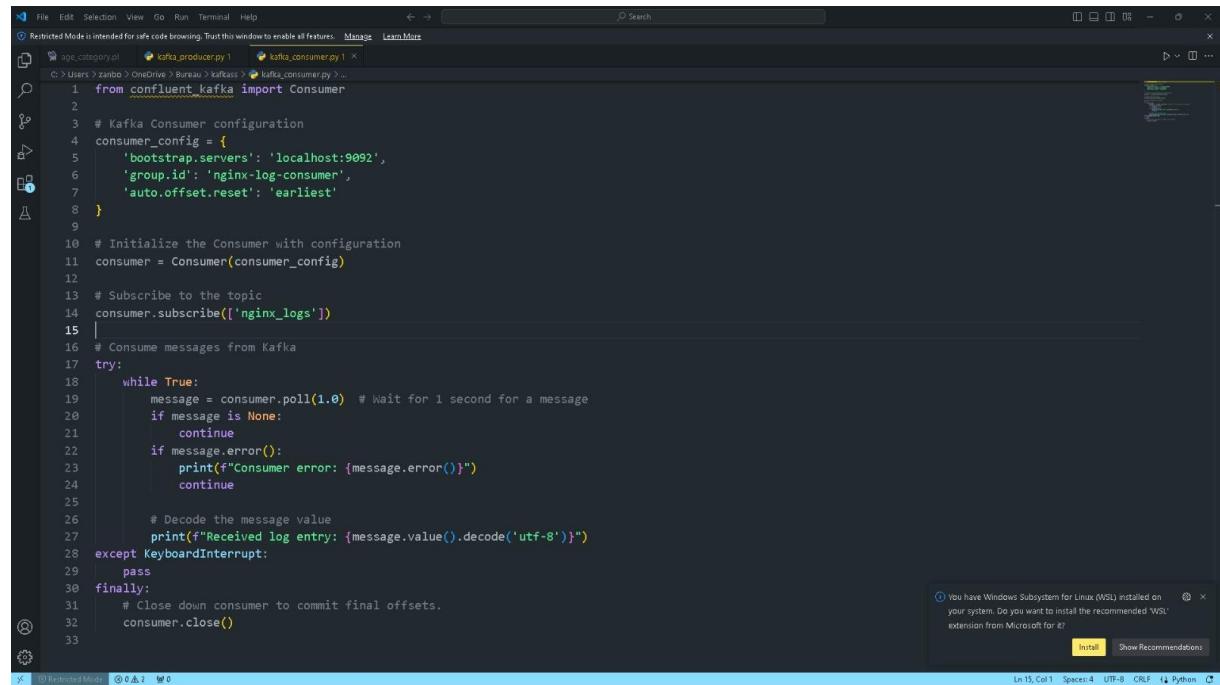
In this part of the assignment, we set up a Python application to send (produce) and receive (consume) log file entries using Kafka. The log files we used are from an Nginx web server. We first created a Python script to read log entries from a file ( a text file )and send them to Kafka. Then, we wrote another script to read these entries from Kafka, mimicking how real-world applications might process log data in real-time.

#### Development Environment and Tools:

**Operating System:** The scripts are run on a Windows 10 machine

**Python Environment:** Python 3 is used

#### Producer and consumer set ups:



```
File Edit Selection View Go Run Terminal Help ⓘ Restricted Mode is intended for safe code browsing. Trust this window to enable all features. Manage Learn More
C:\Users\zanto>OneDrive>Bureau>kafka> kafka_producer.py  kafka_consumer.py 1
C:\Users\zanto>OneDrive>Bureau>kafka> kafka_consumer.py ...
1  from confluent_kafka import Consumer
2
3  # Kafka Consumer configuration
4  consumer_config = {
5      'bootstrap.servers': 'localhost:9092',
6      'group.id': 'nginx-log-consumer',
7      'auto.offset.reset': 'earliest'
8  }
9
10 # Initialize the Consumer with configuration
11 consumer = Consumer(consumer_config)
12
13 # Subscribe to the topic
14 consumer.subscribe(['nginx_logs'])
15 |
16 # Consume messages from Kafka
17 try:
18     while True:
19         message = consumer.poll(1.0) # Wait for 1 second for a message
20         if message is None:
21             continue
22         if message.error():
23             print(f"Consumer error: {message.error()}")
24             continue
25
26         # Decode the message value
27         print(f"Received log entry: {message.value().decode('utf-8')}")
28 except KeyboardInterrupt:
29     pass
30 finally:
31     # Close down consumer to commit final offsets.
32     consumer.close()
33
```

You have Windows Subsystem for Linux (WSL) installed on your system. Do you want to install the recommended 'WSL' extension from Microsoft for it?

Install Show Recommendations

Ln 15, Col 1 Spacing: 4 UTF-8 CRLF Python

The screenshot shows a code editor with two tabs open: `kafka_producer.py` and `kafka_consumer.py`. The `kafka_producer.py` tab is active, displaying the following Python code:

```
age_category.py kafka_producer.py kafka_consumer.py
C:\Users\zando>OneDrive>Bureau>Kafkass>Kafka_producer.py ...
1 #from confluent_kafka import Producer
2 import time
3
4 # Callback function to handle delivery reports
5 def delivery_report(err, msg):
6     if err is not None:
7         print('Message delivery failed: {}'.format(err))
8     else:
9         print('Message delivered to {} [{}].'.format(msg.topic(), msg.partition()))
10
11 # Initialize the Kafka Producer with Confluent Kafka
12 producer = Producer({'bootstrap.servers': 'localhost:9092'})
13
14 # Updated path to the Nginx log file
15 log_file_path = r'C:\Users\zando\OneDrive\Bureau\kafkass\access.txt'
16
17
18 # Stream log entries to Kafka
19 with open(log_file_path, 'r') as file:
20     for line in file:
21         # Send log line to Kafka topic 'nginx_logs' with UTF-8 encoding
22         producer.produce('nginx_logs', value=line.encode('utf-8'), callback=delivery_report)
23         print(f'Sent: {line.strip()}')
24         time.sleep(1) # Add delay for easier demonstration
25
26 # Wait for any outstanding messages to be delivered
27 producer.flush()
28
29
30
```

The status bar at the bottom indicates "Restricted Mode" and "0.0.2 - WSL". A message box in the bottom right corner says: "You have Windows Subsystem for Linux (WSL) installed on your system. Do you want to install the recommended WSL extension from Microsoft for it?" with "Install" and "Show Recommendations" buttons.

## Python Libraries Used

- **Confluent Kafka Python Library:** This is the main library used for creating Kafka producers and consumers
    - **Installation:** Before running the scripts, we installed the Confluent Kafka Python library using pip, Python's package installer, with the command:

*pip install confluent-kafka*

# Kafka and Zookeeper Setup

- **Apache Kafka:** A distributed streaming platform that handles publishing and subscribing to streams of records. It is used here to manage the flow of data from the Nginx logs.
  - **Zookeeper:** Used by Kafka for maintaining configuration information, naming, providing distributed synchronization, and providing group services. All of these facilities help in managing the Kafka cluster.
    - **Docker Containers:** Both Kafka and Zookeeper are run as Docker containers, which isolates their environment and simplifies deployment

## Python Producer Script

- **Purpose:** The producer script reads from a Nginx access log file and sends each line as a message to a Kafka topic called `nginx_logs`.
  - **Key Functions:**
    - **File Handling:** Python's built-in `open()` function is used to read from the Nginx log file.

- **Message Sending:** The produce() method of the Kafka Producer instance is used to send messages to the Kafka server. Messages are encoded in UTF-8 to handle any text data properly.
  - **Callback for Delivery Reports:** This provides asynchronous confirmation of message delivery, helping to debug and ensure reliability in message delivery.

## Python Consumer Script

- **Purpose:** The consumer script listens for messages on the nginx\_logs topic and processes each message as it comes in.
  - **Key Functions:**
    - **Message Polling:** The poll() method is used to wait for new messages from the Kafka server, demonstrating a typical event loop used in real-time data processing applications.
    - **Error Handling:** The script checks for errors in each message, which helps in robustly handling real-world scenarios where network issues or Kafka errors might occur.

**We run the producer first then the consumer:**

## Result:

```
Windows Pow X Windows Pow X Windows Pow X Windows Pow X Command Pn X + - Microsoft Windows [Version 10.0.22621.4317]
(c) Microsoft Corporation. All rights reserved.

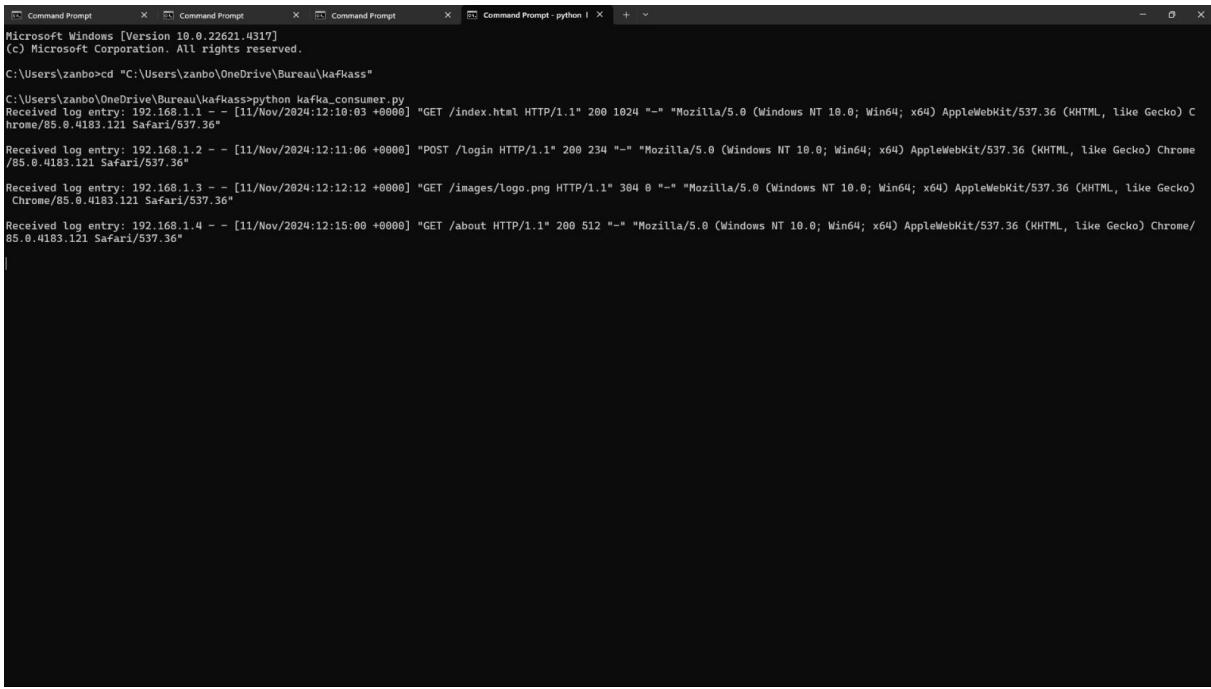
C:\Users\zanko>cd "C:\Users\zanko\OneDrive\Bureau\kafkass"

C:\Users\zanko\OneDrive\Bureau\kafkass>python kafka_producer.py

[11/Nov/2024:12:10:03 +0000] "GET /index.html HTTP/1.1" 200 1024 "-" "Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/85.0.4183.121 Safari/537.36"
Sent: 192.168.1.1 -- [11/Nov/2024:12:10:03 +0000] "GET /index.html HTTP/1.1" 200 1024 "-" "Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/85.0.4183.121 Safari/537.36"
Sent: 192.168.1.2 -- [11/Nov/2024:12:11:06 +0000] "POST /login HTTP/1.1" 200 234 "-" "Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/85.0.4183.121 Safari/537.36"
Sent: 192.168.1.3 -- [11/Nov/2024:12:12:12 +0000] "GET /images/logo.png HTTP/1.1" 304 0 "-" "Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/85.0.4183.121 Safari/537.36"

Sent: 192.168.1.4 -- [11/Nov/2024:12:15:08 +0000] "GET /about HTTP/1.1" 200 512 "-" "Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/85.0.4183.121 Safari/537.36"
Message delivered to nginx_logs [0]

C:\Users\zanko\OneDrive\Bureau\kafkass>
```



The screenshot shows a Windows Command Prompt window with four tabs open. The active tab displays log entries from a Kafka consumer script named `kafka_consumer.py`. The logs show various client requests and their responses, including browser details like Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/85.0.4183.121 Safari/537.36. The log entries are timestamped and include details such as IP address, port, method, URL, status code, and response size.

```
Microsoft Windows [Version 10.0.22621.4317]
(c) Microsoft Corporation. All rights reserved.

C:\Users\zanco>cd "C:\Users\zanco\OneDrive\Bureau\kafkass"

C:\Users\zanco>python kafka_consumer.py
Received log entry: 192.168.1.1 -- [11/Nov/2024:12:10:03 +0000] "GET /index.html HTTP/1.1" 200 1024 "-" "Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/85.0.4183.121 Safari/537.36"
Received log entry: 192.168.1.2 -- [11/Nov/2024:12:11:06 +0000] "POST /login HTTP/1.1" 200 234 "-" "Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/85.0.4183.121 Safari/537.36"
Received log entry: 192.168.1.3 -- [11/Nov/2024:12:12:12 +0000] "GET /images/logo.png HTTP/1.1" 304 0 "-" "Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/85.0.4183.121 Safari/537.36"
Received log entry: 192.168.1.4 -- [11/Nov/2024:12:15:00 +0000] "GET /about HTTP/1.1" 200 512 "-" "Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/85.0.4183.121 Safari/537.36"
```

Each line in the Command Prompt window begins with "Received log entry:", indicating that the consumer script is successfully retrieving messages. These messages are log entries originally written to an Nginx access log file and then published to Kafka by the producer script. The details shown in each log entry include the client's IP address, the timestamp of the request, the HTTP method and resource requested, the HTTP status code and response size, and the client's browser information. For instance, entries like "GET /index.html HTTP/1.1" followed by "200 1024" denote a successful fetch of the 'index.html' page with a 200 OK status and a response size of 1024 bytes.

This output not only validates the correct functioning of the consumer script but also showcases real-time data processing capabilities. It demonstrates the integration between the producer and consumer components, affirming that they can communicate effectively through Kafka. Moreover, this visualization of log data is crucial for monitoring the data flow, assisting in debugging and operational oversight, ensuring all parts of the Kafka setup operate cohesively.

## Question 5

In this bonus question, we explored advanced capabilities of Kafka consumer groups by setting them up in two distinct configurations: fanout and load balancing. These configurations allow Kafka to distribute and manage data streams more effectively across multiple consumers, enhancing both the efficiency and reliability of data processing.

## Consumer Groups Concept

- **Consumer Groups:** A consumer group in Kafka consists of multiple consumers that jointly process the data contained in a Kafka topic. Each consumer within a group reads from exclusive partitions of the topic, ensuring that each message is processed just once by the group as a whole.
- **Fanout:** This setup involves multiple consumer groups consuming the same data independently of each other. It's useful for scenarios where different systems or applications need to process the same data in parallel without interfering with each other.
- **Load Balancing:** In this configuration, all consumers belong to the same group but share the work of consuming messages. Kafka automatically distributes the messages among available consumers in the group, balancing the load and increasing throughput.

## Implementation of the work

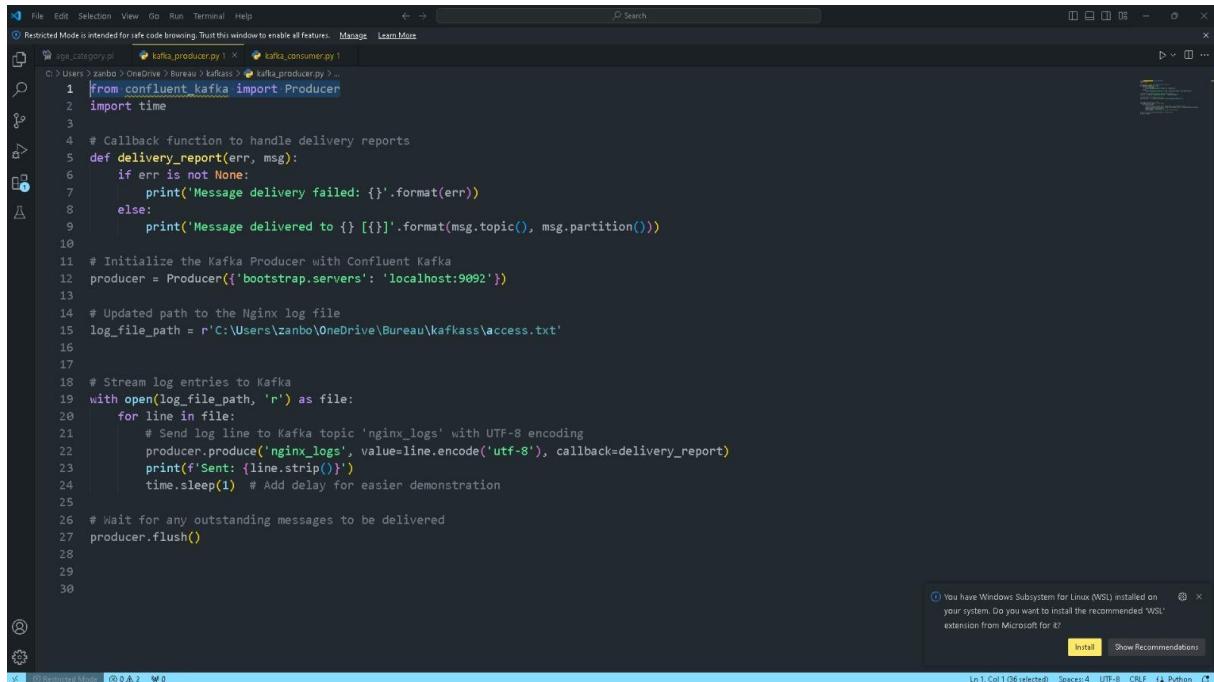
- **Kafka Producer Script:** As with earlier tasks, the producer script sends log entries from an Nginx access log file to a Kafka topic named 'nginx\_logs'. The producer is configured to connect to Kafka running on localhost:9092 and sends messages with UTF-8 encoding to ensure accurate text representation.
- **Fanout Configuration:** We set up multiple consumer groups, each configured to consume messages from the 'nginx\_logs' topic independently. This demonstrates the fanout configuration, where each group processes the same set of data for different purposes or in different ways.
- **Load Balancing Configuration:** We configured multiple consumers under a single group ID to demonstrate load balancing. Kafka distributes the incoming messages among all consumers in this group, effectively balancing the workload and enhancing processing speed.

## Kafka Producer Script

- **Content:** Shows the Kafka producer script in an IDE. This script is responsible for reading Nginx log entries from a file and sending them to a Kafka topic (nginx\_logs).
- **Highlights:**
  - Initialization of the Kafka producer with the server address localhost:9092.
  - Sending log entries with a callback to confirm delivery.
  - This script is crucial for supplying data to Kafka, which is then consumed by different consumer groups.

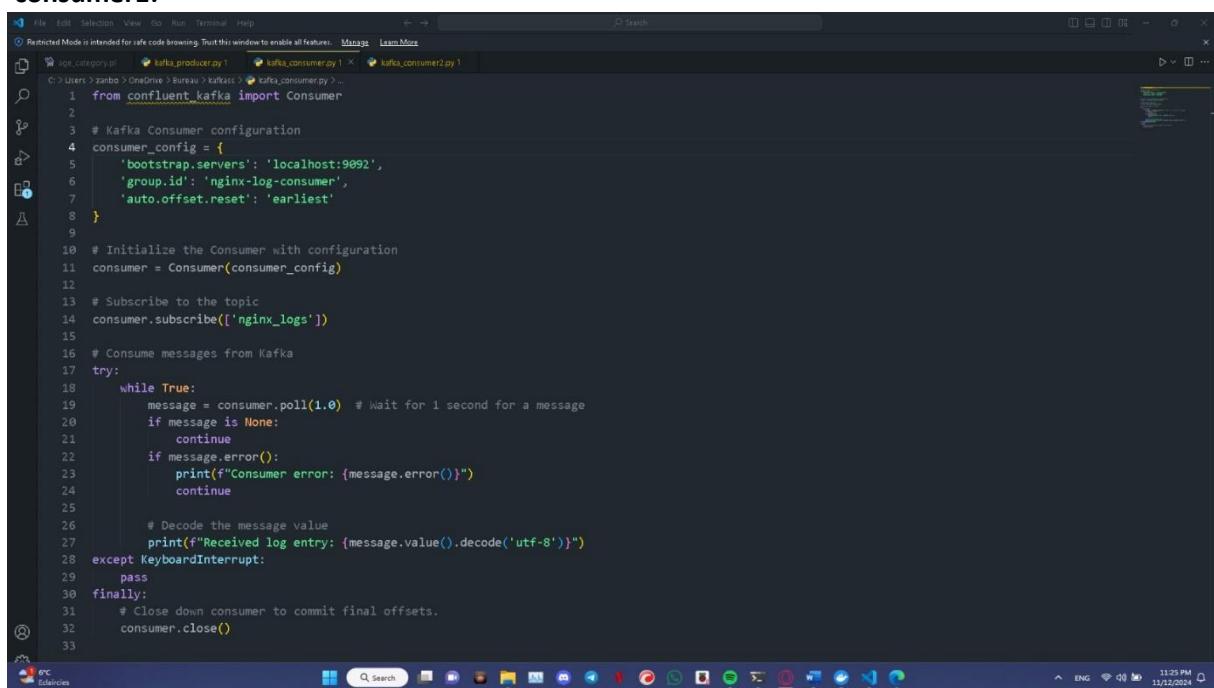
Here is its implementation : fig 1

Producer:



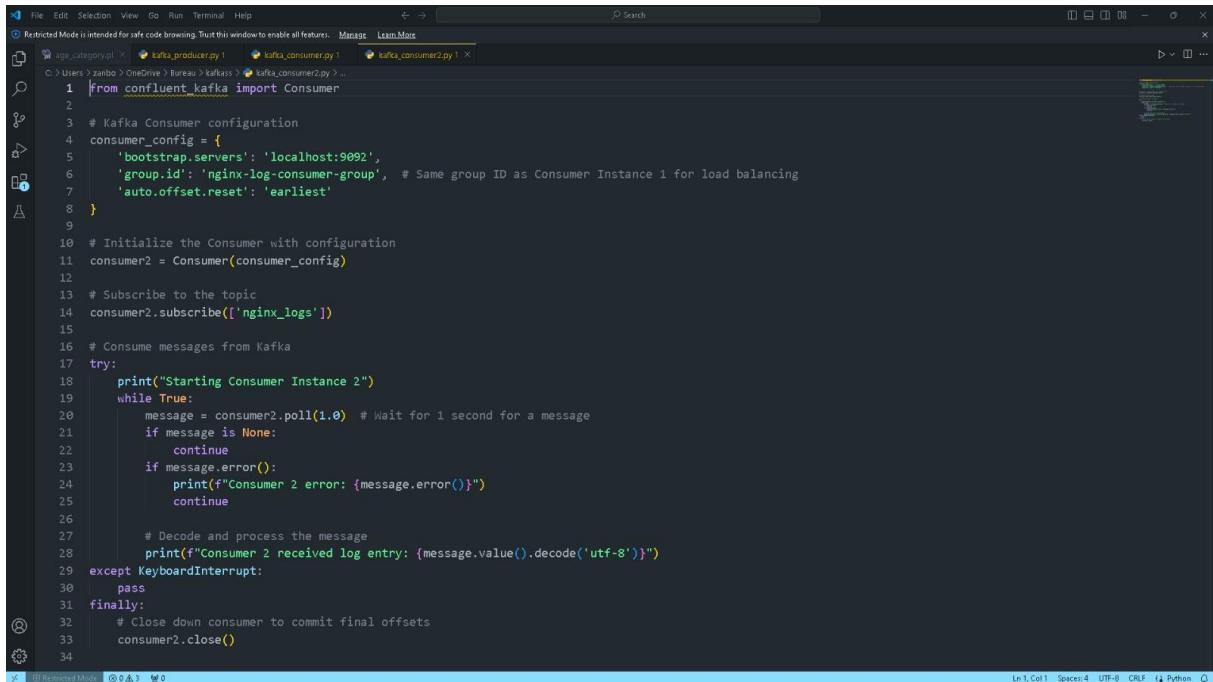
```
File Edit Selection View Go Run Terminal Help
Restricted Mode is intended for safe code browsing. Trust this window to enable all features. Manage Learn More
C:\Users\zando>OneDrive>Bureau>kafkass>kafka_producer.py ...
1 #from confluent_kafka import Producer
2 import time
3
4 # Callback function to handle delivery reports
5 def delivery_report(err, msg):
6     if err is not None:
7         print('Message delivery failed: {}'.format(err))
8     else:
9         print('Message delivered to {} [{}].format(msg.topic(), msg.partition())')
10
11 # Initialize the Kafka Producer with Confluent Kafka
12 producer = Producer({'bootstrap.servers': 'localhost:9092'})
13
14 # Updated path to the Nginx log file
15 log_file_path = r'C:\Users\zando\OneDrive\Bureau\kafkass\access.txt'
16
17
18 # Stream log entries to Kafka
19 with open(log_file_path, 'r') as file:
20     for line in file:
21         # Send log line to Kafka topic 'nginx_logs' with UTF-8 encoding
22         producer.produce('nginx_logs', value=line.encode('utf-8'), callback=delivery_report)
23         print(f'Sent: {line.strip()}')
24         time.sleep(1) # Add delay for easier demonstration
25
26 # Wait for any outstanding messages to be delivered
27 producer.flush()
28
29
30
```

### consumer1:



```
File Edit Selection View Go Run Terminal Help
Restricted Mode is intended for safe code browsing. Trust this window to enable all features. Manage Learn More
C:\Users\zando>OneDrive>Bureau>kafkass>kafka_producer.py ...
1 from confluent_kafka import Consumer
2
3 # Kafka Consumer configuration
4 consumer_config = {
5     'bootstrap.servers': 'localhost:9092',
6     'group.id': 'nginx-log-consumer',
7     'auto.offset.reset': 'earliest'
8 }
9
10 # Initialize the Consumer with configuration
11 consumer = Consumer(consumer_config)
12
13 # Subscribe to the topic
14 consumer.subscribe(['nginx_logs'])
15
16 # Consume messages from Kafka
17 try:
18     while True:
19         message = consumer.poll(1.0) # Wait for 1 second for a message
20         if message is None:
21             continue
22         if message.error():
23             print("Consumer error: {}".format(message.error()))
24             continue
25
26         # Decode the message value
27         print("Received log entry: {}".format(message.value().decode('utf-8')))
28 except KeyboardInterrupt:
29     pass
30 finally:
31     # Close down consumer to commit final offsets.
32     consumer.close()
```

### consumer2:

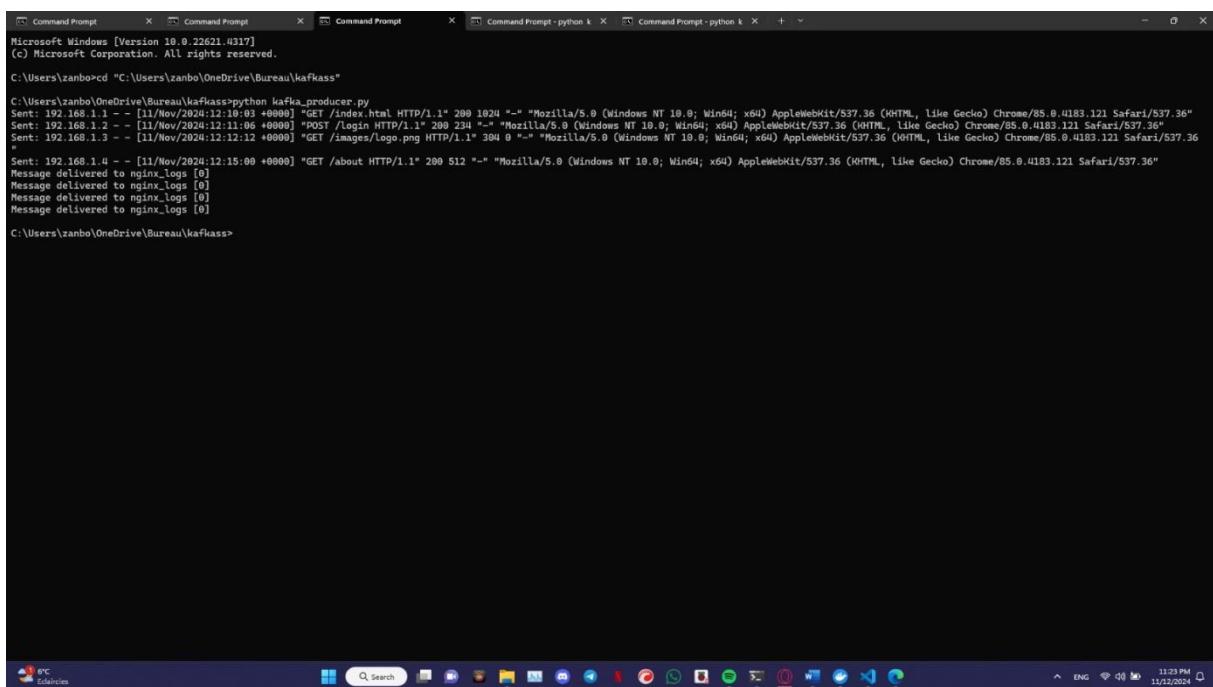


A screenshot of a code editor window titled "age\_category.py". The code is a Python script for a Kafka consumer. It imports the Consumer class from confluent\_kafka. The configuration consumer\_config specifies bootstrap.servers as 'localhost:9092', group.id as 'nginx-log-consumer-group', and auto.offset.reset as 'earliest'. The consumer is initialized with this configuration, subscribed to the topic 'nginx\_logs', and then enters a loop to poll for messages every 1 second. If a message is received, it is printed. The script handles KeyboardInterrupt exceptions and closes the consumer when it exits.

```
1 #from confluent_kafka import Consumer
2
3 # Kafka Consumer configuration
4 consumer_config = {
5     'bootstrap.servers': 'localhost:9092',
6     'group.id': 'nginx-log-consumer-group', # Same group ID as Consumer Instance 1 for load balancing
7     'auto.offset.reset': 'earliest'
8 }
9
10 # Initialize the Consumer with configuration
11 consumer2 = Consumer(consumer_config)
12
13 # Subscribe to the topic
14 consumer2.subscribe(['nginx_logs'])
15
16 # Consume messages from Kafka
17 try:
18     print("Starting Consumer Instance 2")
19     while True:
20         message = consumer2.poll(1.0) # Wait for 1 second for a message
21         if message is None:
22             continue
23         if message.error():
24             print(f"Consumer 2 error: {message.error()}")
25             continue
26
27         # Decode and process the message
28         print(f"Consumer 2 received log entry: {message.value().decode('utf-8')}")
29 except KeyboardInterrupt:
30     pass
31 finally:
32     # Close down consumer to commit final offsets
33     consumer2.close()
34
```

Then the results :

### Running the producer first:



A screenshot of a terminal window titled "Command Prompt" showing the output of a Kafka producer script. The command "cd C:\Users\zanco\OneDrive\Bureau\kafkass" is run, followed by "python kafka\_producer.py". The producer sends four messages to the 'nginx\_logs' topic. The first two messages are GET requests for index.html and /login respectively, both from IP 192.168.1.1. The last two messages are GET requests for images/logo.png, also from IP 192.168.1.1. All messages are sent at approximately 12:18:03 on November 11, 2024.

```
C:\Users\zanco\OneDrive\Bureau\kafkass>cd "C:\Users\zanco\OneDrive\Bureau\kafkass"
C:\Users\zanco\OneDrive\Bureau\kafkass>python kafka_producer.py
[2024-11-11 12:18:03,000] {producer.py:40} INFO - Sent: 192.168.1.1 - - [11/Nov/2024:12:18:03 +0000] "GET /index.html HTTP/1.1" 200 1824 "-" "Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/85.0.4183.121 Safari/537.36"
[2024-11-11 12:18:03,000] {producer.py:40} INFO - Sent: 192.168.1.2 - - [11/Nov/2024:12:18:03 +0000] "POST /Login HTTP/1.1" 200 234 "-" "Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/85.0.4183.121 Safari/537.36"
[2024-11-11 12:18:03,000] {producer.py:40} INFO - Sent: 192.168.1.3 - - [11/Nov/2024:12:18:03 +0000] "GET /images/logo.png HTTP/1.1" 200 394 "-" "Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/85.0.4183.121 Safari/537.36"
[2024-11-11 12:18:03,000] {producer.py:40} INFO - Sent: 192.168.1.4 - - [11/Nov/2024:12:18:03 +0000] "GET /about HTTP/1.1" 200 512 "-" "Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/85.0.4183.121 Safari/537.36"
Message delivered to nginx_logs [0]
```

Then running the consumers:

```
Microsoft Windows [Version 10.0.22621.4317]
(c) Microsoft Corporation. All rights reserved.

C:\Users\zanko>cd "C:\Users\zanko\OneDrive\Bureau\kafkass"

C:\Users\zanko\OneDrive\Bureau\kafkass>python kafka_consumer2.py
Starting Consumer Instance 2
Consumer 2 received log entry: 192.168.1.1 -- [11/Nov/2024:12:10:03 +0000] "GET /index.html HTTP/1.1" 200 1024 "-" "Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/85.0.4183.121 Safari/537.36"

Consumer 2 received log entry: 192.168.1.2 -- [11/Nov/2024:12:11:06 +0000] "POST /login HTTP/1.1" 200 234 "-" "Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/85.0.4183.121 Safari/537.36"

Consumer 2 received log entry: 192.168.1.3 -- [11/Nov/2024:12:12:12 +0000] "GET /images/Logo.png HTTP/1.1" 304 0 "-" "Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/85.0.4183.121 Safari/537.36"

Consumer 2 received log entry: 192.168.1.4 -- [11/Nov/2024:12:15:00 +0000] "GET /about HTTP/1.1" 200 512 "-" "Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/85.0.4183.121 Safari/537.36"
```

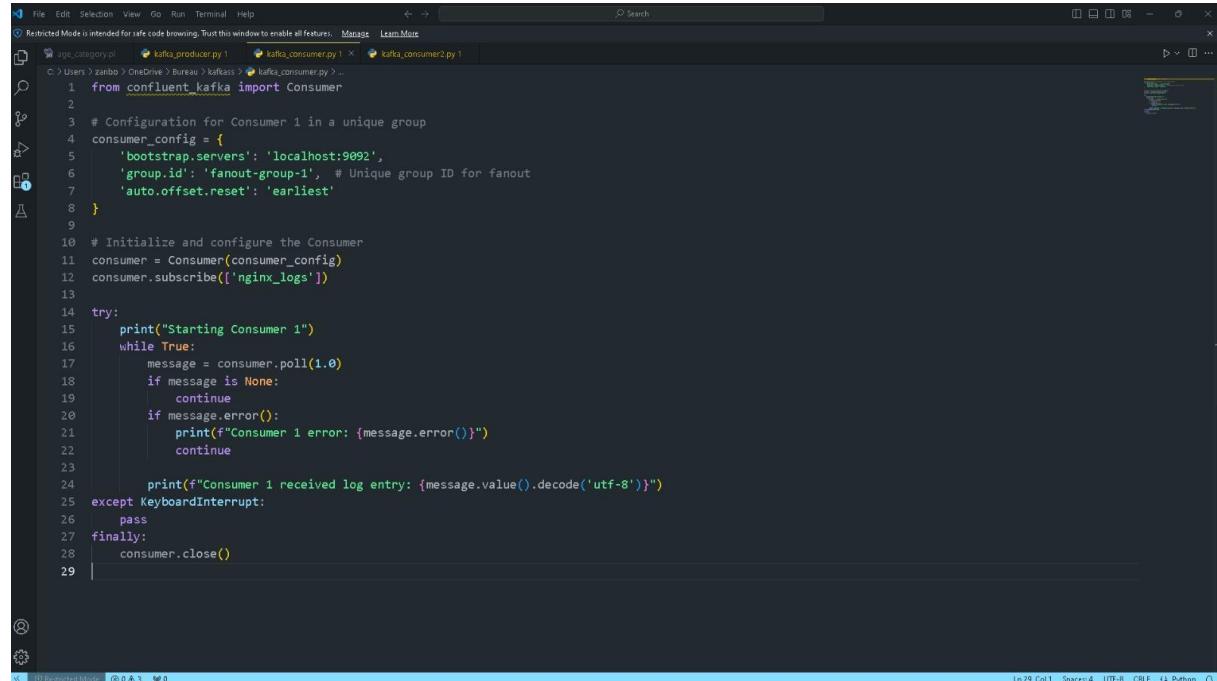
So we see that the consumers received different log entries from the same topic. By setting up multiple consumers within the same group, Kafka's load balancing capability was effectively utilized. This configuration enables scalable data processing, making it well-suited for high-throughput applications that demand efficient message management across distributed systems. Kafka's consumer group load balancing provides a strong solution for applications requiring scalable and reliable data streaming.

## FIG 2 : Kafka Consumer Script for Fanout Configuration

- Displays a Kafka consumer script configured for a fanout scenario. This script is set to consume messages from the `nginx_logs` topic independently from other consumers.

- Consumer is set up with a unique group ID (fanout-group-2), allowing it to consume messages in parallel with others.
- Script includes real-time message consumption logs, showing each entry as it is processed.

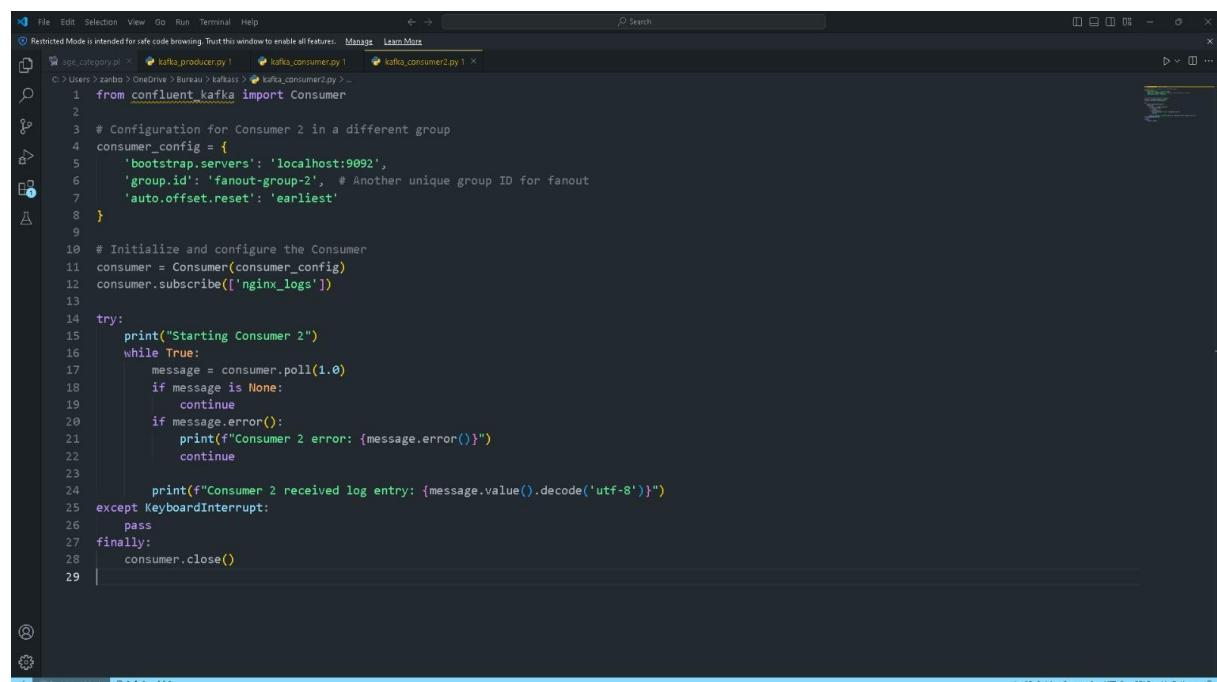
Here is the code for the consumers for the fanout configuration:



```

File Edit Selection View Go Run Terminal Help
Restricted Mode is intended for safe code browsing. Trust this window to enable all features. Manage Learn More
C:\Users\zarbo>OneDrive>Bureau>kafka> kafka_consumer.py ...
1 from confluent_kafka import Consumer
2
3 # Configuration for Consumer 1 in a unique group
4 consumer_config = {
5     'bootstrap.servers': 'localhost:9092',
6     'group.id': 'fanout-group-1', # Unique group ID for fanout
7     'auto.offset.reset': 'earliest'
8 }
9
10 # Initialize and configure the Consumer
11 consumer = Consumer(consumer_config)
12 consumer.subscribe(['nginx_logs'])
13
14 try:
15     print("Starting Consumer 1")
16     while True:
17         message = consumer.poll(1.0)
18         if message is None:
19             continue
20         if message.error():
21             print(f"Consumer 1 error: {message.error()}")
22             continue
23
24         print(f"Consumer 1 received log entry: {message.value().decode('utf-8')}")
25 except KeyboardInterrupt:
26     pass
27 finally:
28     consumer.close()
29

```



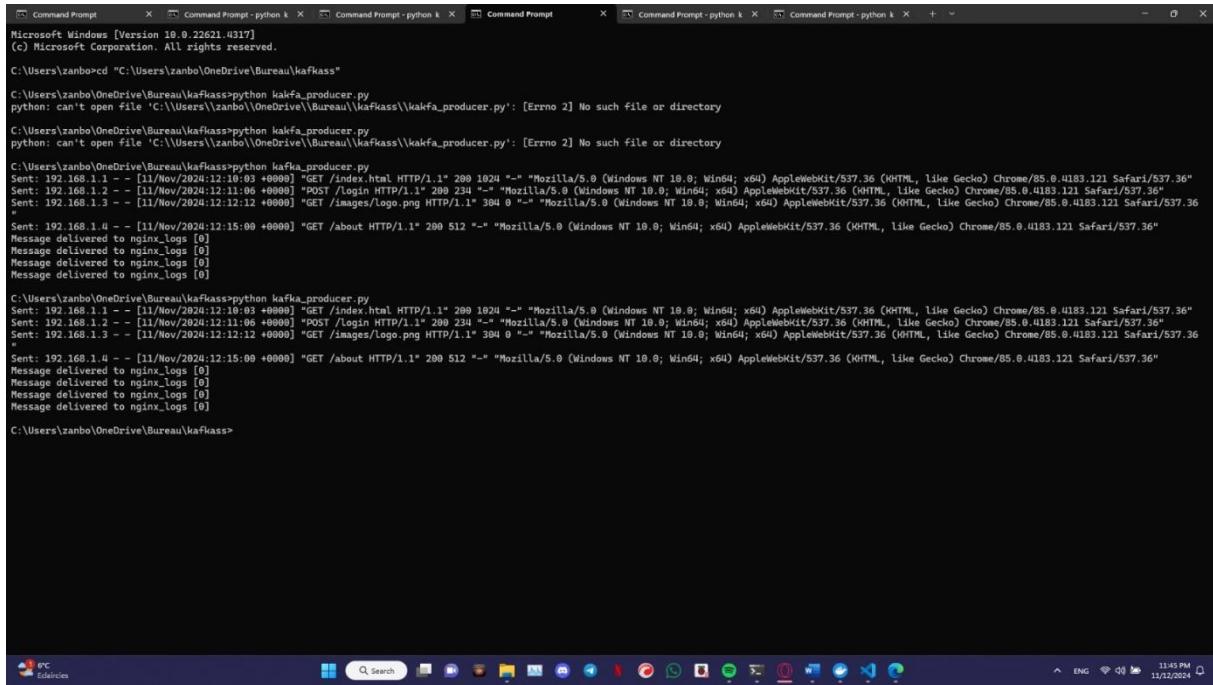
```

File Edit Selection View Go Run Terminal Help
Restricted Mode is intended for safe code browsing. Trust this window to enable all features. Manage Learn More
C:\Users\zarbo>OneDrive>Bureau>kafka> kafka_consumer2.py ...
1 from confluent_kafka import Consumer
2
3 # Configuration for Consumer 2 in a different group
4 consumer_config = {
5     'bootstrap.servers': 'localhost:9092',
6     'group.id': 'fanout-group-2', # Another unique group ID for fanout
7     'auto.offset.reset': 'earliest'
8 }
9
10 # Initialize and configure the Consumer
11 consumer = Consumer(consumer_config)
12 consumer.subscribe(['nginx_logs'])
13
14 try:
15     print("Starting Consumer 2")
16     while True:
17         message = consumer.poll(1.0)
18         if message is None:
19             continue
20         if message.error():
21             print(f"Consumer 2 error: {message.error()}")
22             continue
23
24         print(f"Consumer 2 received log entry: {message.value().decode('utf-8')}")
25 except KeyboardInterrupt:
26     pass
27 finally:
28     consumer.close()
29

```

**Fig 3 : Messages are successfully sent to the Kafka topic, as indicated by "Message delivered to nginx\_logs [0]".**

- Shows the real-time logging of HTTP requests being processed and confirmed as sent.
- Validates that the producer script is actively interacting with Kafka and the messages are queued correctly in the topic.



```
Microsoft Windows [Version 10.0.22621.4317]
(c) Microsoft Corporation. All rights reserved.

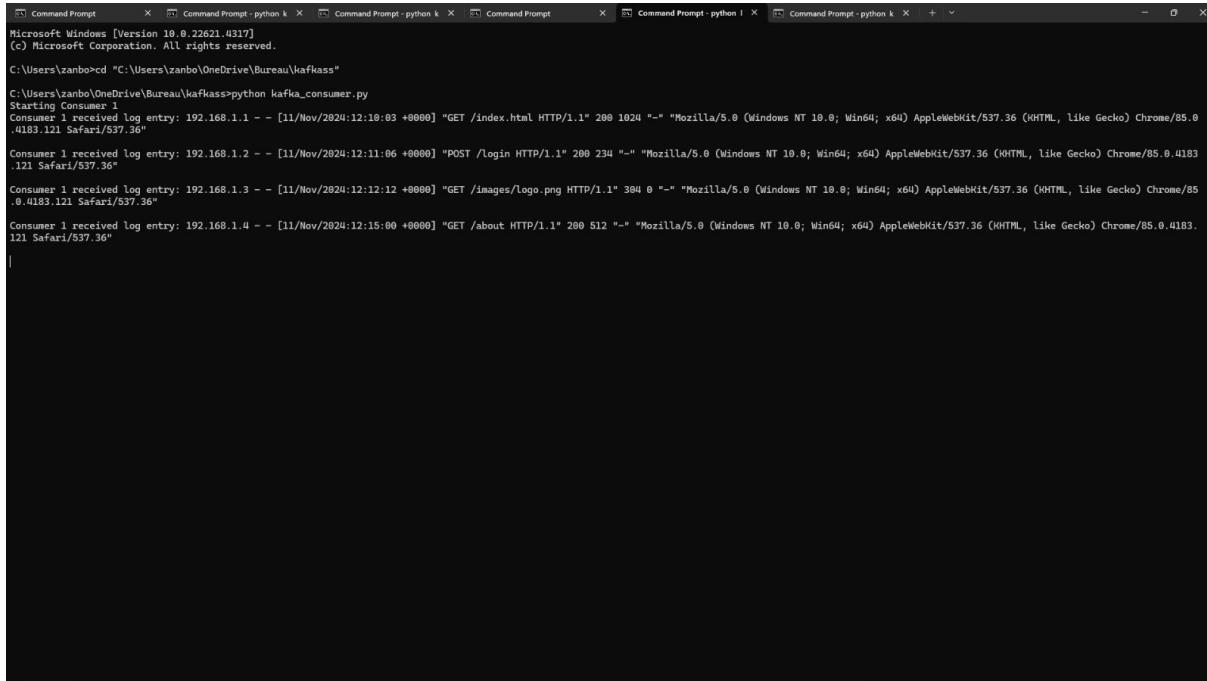
C:\Users\zanco>cd "C:\Users\zanco\OneDrive\Bureau\kafkass"
C:\Users\zanco>python kafka_producer.py
python: can't open file 'C:\\\\users\\\\zanco\\\\OneDrive\\\\Bureau\\\\kafkass\\\\kafka_producer.py': [Errno 2] No such file or directory

C:\Users\zanco>python kafka_producer.py
python: can't open file 'C:\\\\users\\\\zanco\\\\OneDrive\\\\Bureau\\\\kafkass\\\\kafka_producer.py': [Errno 2] No such file or directory

C:\Users\zanco>python kafka_producer.py
Sent: 192.168.1.1 - - [11/Nov/2024:12:11:08 +0000] "GET /index.html HTTP/1.1" 200 1024 "-" "Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/85.0.4183.121 Safari/537.36"
Sent: 192.168.1.2 - - [11/Nov/2024:12:11:08 +0000] "POST /login HTTP/1.1" 200 234 "-" "Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/85.0.4183.121 Safari/537.36"
Sent: 192.168.1.3 - - [11/Nov/2024:12:12:12 +0000] "GET /images/logo.png HTTP/1.1" 304 0 "-" "Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/85.0.4183.121 Safari/537.36"
Sent: 192.168.1.4 - - [11/Nov/2024:12:15:08 +0000] "GET /about HTTP/1.1" 200 512 "-" "Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/85.0.4183.121 Safari/537.36"
Message delivered to nginx_logs [0]

C:\Users\zanco>python kafka_producer.py
Sent: 192.168.1.1 - - [11/Nov/2024:12:18:03 +0000] "GET /index.html HTTP/1.1" 200 1024 "-" "Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/85.0.4183.121 Safari/537.36"
Sent: 192.168.1.2 - - [11/Nov/2024:12:11:06 +0000] "POST /Login HTTP/1.1" 200 234 "-" "Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/85.0.4183.121 Safari/537.36"
Sent: 192.168.1.3 - - [11/Nov/2024:12:12:12 +0000] "GET /images/logo.png HTTP/1.1" 304 0 "-" "Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/85.0.4183.121 Safari/537.36"
"
Sent: 192.168.1.4 - - [11/Nov/2024:12:15:00 +0000] "GET /about HTTP/1.1" 200 512 "-" "Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/85.0.4183.121 Safari/537.36"
Message delivered to nginx_logs [0]
```

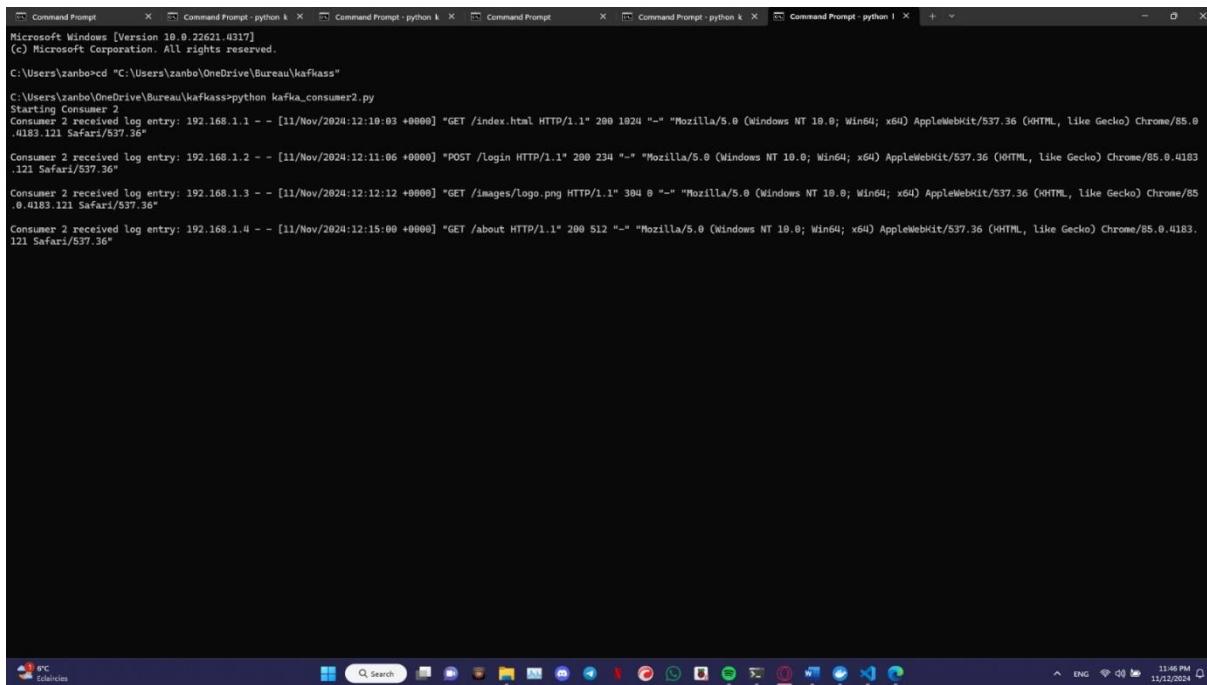
**Fig 4 :** Displays multiple log entries being received and processed by the consumer.



```
Microsoft Windows [Version 10.0.22621.4317]
(c) Microsoft Corporation. All rights reserved.

C:\Users\zanco>cd "C:\Users\zanco\OneDrive\Bureau\kafkass"
C:\Users\zanco\OneDrive\Bureau\kafkass>python kafka_consumer.py
Starting Consumer 1
Consumer 1 received log entry: 192.168.1.1 - - [11/Nov/2024:12:10:03 +0000] "GET /index.html HTTP/1.1" 200 1024 "-" "Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/85.0.4183.121 Safari/537.36"
Consumer 1 received log entry: 192.168.1.2 - - [11/Nov/2024:12:11:06 +0000] "POST /login HTTP/1.1" 200 234 "-" "Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/85.0.4183.121 Safari/537.36"
Consumer 1 received log entry: 192.168.1.3 - - [11/Nov/2024:12:12:12 +0000] "GET /images/logo.png HTTP/1.1" 304 0 "-" "Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/85.0.4183.121 Safari/537.36"
Consumer 1 received log entry: 192.168.1.4 - - [11/Nov/2024:12:15:00 +0000] "GET /about HTTP/1.1" 200 512 "-" "Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/85.0.4183.121 Safari/537.36"
```

**FIG 5 : Each consumer receives and logs the same data, demonstrating the fanout model where each consumer group receives all messages independently. Useful for scenarios where multiple applications or services need to process the same data simultaneously without interference.**



```
Microsoft Windows [Version 10.0.22621.4317]
(c) Microsoft Corporation. All rights reserved.

C:\Users\zanco>cd "C:\Users\zanco\OneDrive\Bureau\kafkass"
C:\Users\zanco\OneDrive\Bureau\kafkass>python kafka_consumer2.py
Starting Consumer 2
Consumer 2 received log entry: 192.168.1.1 - - [11/Nov/2024:12:10:03 +0000] "GET /index.html HTTP/1.1" 200 1024 "-" "Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/85.0.4183.121 Safari/537.36"
Consumer 2 received log entry: 192.168.1.2 - - [11/Nov/2024:12:11:06 +0000] "POST /login HTTP/1.1" 200 234 "-" "Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/85.0.4183.121 Safari/537.36"
Consumer 2 received log entry: 192.168.1.3 - - [11/Nov/2024:12:12:12 +0000] "GET /images/logo.png HTTP/1.1" 304 0 "-" "Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/85.0.4183.121 Safari/537.36"
Consumer 2 received log entry: 192.168.1.4 - - [11/Nov/2024:12:15:00 +0000] "GET /about HTTP/1.1" 200 512 "-" "Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/85.0.4183.121 Safari/537.36"
```

**So here for the fanout we see that the consumers receive the same log entries. The configuration show that the each one act independently, receiving a full copy of the log.**

### References

1. Kreps, J., Narkhede, N., & Rao, J. (2011). *Kafka: A Distributed Messaging System for Log Processing*. LinkedIn. Available at: <https://kafka.apache.org>
2. Confluent. (2023). *Apache Kafka Overview*. Confluent Documentation. Available at: <https://docs.confluent.io>
3. Amazon Web Services. (2023). *What is Amazon Simple Notification Service (SNS)*? Available at: <https://aws.amazon.com/sns/>
4. Pivotal Software. (2023). *RabbitMQ Overview*. RabbitMQ Documentation. Available at: <https://www.rabbitmq.com>
5. Red Hat. (2023). *Apache ActiveMQ Overview*. Available at: <https://activemq.apache.org>
6. Kafka Documentation. *Producer and Consumer API*. Available at: <https://kafka.apache.org/documentation>